

Five signs your AI is *lying* to you

Run this in the next twenty minutes, on your own stack, before you trust the next number it hands you.

TRACE EACH ONE AGAINST YOUR OWN STACK

Check the box only after you have actually traced it, not after you have thought about it.

SIGNAL & WHAT IT LOOKS LIKE	SELF-CHECK
☐ ¹ Model does the arithmetic. <i>Model outputs a total, score, or forecast. If a model computes it, you have a silent-error generator with no seam between reading and computing.</i>	<i>Trace one number in your report: model, or code?</i>
☐ ² A decision you can't explain. <i>Only answer is "the model judged it." You cannot defend a decision, or debug it, with no rule to point at.</i>	<i>Name the exact rule behind one automated decision.</i>
☐ ³ A green check that never fails. <i>Your top test has never once failed. It may be measuring something adjacent to what you actually care about.</i>	<i>When did that check last actually fail? Never?</i>
☐ ⁴ A finding with no citation. <i>A claim ships with no source line attached. Treat it as a hallucination until proven otherwise, grounding is not a nicety.</i>	<i>Pick one AI claim: does it name where it came from?</i>
☐ ⁵ Shipped and never re-read. <i>Finished output nobody has reopened since. Finished reads as correct, it usually isn't re-checked, which is exactly where errors hide.</i>	<i>Name one deliverable this quarter you never reopened.</i>

Self-tally: how many of the five did you just fail? Zero is rare. One is normal. Three or more means something in your stack is quietly wrong right now, not hypothetically.

Why this matters more than it feels like

None of these five require an AI expert to catch. They require someone willing to trace one number, one decision, or one claim back to where it actually came from, instead of trusting the output because it looks finished.

What a passing stack looks like

The arithmetic runs in code, not the model. Every automated decision has a rule you can point to. The tests that matter have failed at least once. Every claim carries a source. Every deliverable gets reopened.

NEXT STEP

If you scored two or more, that is not bad luck, it is a pattern I have already found once in someone else's pipeline this month.
Book a free 20-minute brief and bring your worst one, we will trace it live.
Want more of this first: **connect on LinkedIn**.