

THE APPROVAL GATE TEST

AUG 2026

Your approval step might just be a *label*

A two-minute test to find out whether anything in your agent workflow can actually say no. Most cannot.

Found by testing a set of published agent-pattern examples that call a step an "approval gate," then checking whether anything in the code actually blocks the action if that label says no. Most of the time, nothing does.

● THE TRIGGER

You added a step that says "route to approval" or "needs human review" somewhere in an agent or automation you built. You assumed that step was a safety net. It might be. It might also be a field in a JSON output that nothing downstream ever reads.

● THE INSIGHT

Across a set of published agent-pattern examples, one pattern kept showing up under the name "approval gate": a router labels an input with an approval tier, and the demo stops there. Nothing in the code actually blocks the action if that label says "needs approval." The label is the whole feature. A real gate is different: it is a check that runs outside the model's own judgment, and the action cannot proceed until that check passes.

A label

A router marks an input with an approval tier. The demo stops there. **Nothing in the code blocks the action** if that label says "needs approval."

A gate

A check runs **outside the model's own judgment**. The action cannot proceed until that check passes, no matter how confident the model's own output looked.

● THE TEST *run it today, on your own workflow*

1. **Find** the approval or review step in your own workflow.
2. **Ask** one question: if that step said "no," what line of code, or what script, would actually stop the action from happening?
3. **Point** to it. Not describe it. Point to it.
4. If the honest answer is "the model decided" or "there isn't one," **you have a label, not a gate**.

Most workflows fail step 3. That is not a verdict on the workflow, it is the finding the rest of this test exists to surface.

● WHAT A REAL GATE LOOKS LIKE

A gate is deterministic and external to the model. One concrete shape, from a content pipeline I run day to day: nothing gets marked ready to ship until a human-set review status says so, and separately, two small scripts have to exit clean, pass or fail, no judgment call, before the ship step is allowed to run. The model that drafted the content never grades whether its own gate passed. That is what makes it a gate instead of a label: the block is enforced somewhere the model cannot talk its way around.

● WHAT YOU NOW KNOW

You either found a real gate, rare, and worth confirming twice, or you found a label wearing a gate's name, common. Either way, you now know which one you are running, and that is the finding an AI Workflow Audit exists to surface, just applied to one step instead of your whole stack.

If you found a gate

Rare. Confirm it twice: check that the block is a script or a status check that runs outside the model, not a status field that only looks reassuring.

If you found a label

Common. The label is not a failure, it is a starting point: the same finding an AI Workflow Audit exists to surface, just applied to one step instead of the whole stack.

NEXT STEP

If that test turned up a label instead of a gate, it is worth finding out what else in the same workflow is a label. **Book a free 20-minute brief** and we will run the same test on the rest of your pipeline.

Want more of this first: **connect on LinkedIn**.